

Step-by-Step Beginner Guide: Deploying a Zero-Cost ML API on Render

Prerequisites: - GitHub account - Render account (free) - Python 3.8+ locally, pip, and optional Docker

Day 1: Local setup

1. Clone the repo. Run `python -m pip install -r requirements.txt`.
2. Run `python train.py` to create model.joblib.
3. Run `uvicorn app:app --reload --host 0.0.0.0 --port 8000` and test locally using curl or Postman.

Day 2: Push to GitHub & Deploy on Render

1. Create a GitHub repo and push all files.
2. On Render, create a new Web Service and connect to the GitHub repo.
3. Set Build Command: `pip install -r requirements.txt`
4. Set Start Command: `uvicorn app:app --host 0.0.0.0 --port 10000`
5. Select Free instance. Deploy and wait for the live URL.

Day 3: Validate, Test & Report

1. Use curl to run multiple concurrent requests and observe Render logs.
2. Capture screenshots of the Render dashboard and logs showing cold-starts or request spikes.
3. Prepare submission: one-page proposal, architecture PNG, and final report docx.